

# Linear Algebra Optimization in C using OpenMPI

Aidan Levy

Dr. Spickler

May 12th, 2025

# 1 Introduction

In today's data driven world, there is a constant demand for faster and most efficient computing methods. Traditional sequential computing, where tasks are executed one after another by a singular processor, often falls short when tackling complex and large scale problems. While this model remains useful for certain tasks and small scale applications, it fails to meet the performance needs of large scale computations in modern research, analytics, and simulation based fields.

Parallel computing addresses this limitation by distributing workload across multiple processors or cores, enabling simultaneous execution of sub tasks. This approach significantly reduces execution time and increases efficiency, especially for complex mathematical operations and data heavy computations. Parallelism has become integral in many different cases— from national defense and physics simulations to artificial intelligence and high frequency trading— where performance constraints dictate the outcomes of solutions (Barney & Frederick). The Open Message Passing Interface (OpenMPI) implementation has become a core element of solving this problem by facilitating communication and coordination among distributed processes. OpenMPI is widely used in both academic and industry settings due to its scalability, ease of use, and support for high performance clusters.

This research project investigates the use of OpenMPI to optimize linear algebra algorithms written in C, which is a very low level language. Linear algebra remains a central pillar of computer science, with applications ranging from computer graphics and machine learning to economic modeling and engineering. One historic example includes Professor Leontief's use of linear equations in 1949 to model the structure of the United States economy, which demonstrated the importance of these systems (Lay et al., 2020). In this study, emphasis is placed on the efficient execution of operations such as matrix multiplication, row echelon form (REF), reduced row echelon form (RREF), and LU decomposition, which serve as critical tools in solving systems of linear equations. These algorithms are implemented in both sequential and parallel forms to allow for direct performance comparisons. Timing benchmarks collected during testing reveal the extent to which parallelism accelerates execution and under what conditions those gains are most significant. These performance results will be discussed in detail in later sections.

In addition to algorithm optimization, the project involved constructing a custom parallel computing environment. The process of configuring and deploying a multi node OpenMPI cluster provided valuable learning experiences into the logistics and technical requirement of distributed computing. This included configuration of operating systems, secure shell (SSH) communication setup between nodes, implementation of shared storage, and fine tuning of the scripts. The resulting cluster not only enabled controlled benchmarking of parallel code but also served as a replicable model for lightweight high performance computing environments.

## 2 Terminology and MPI Overview

To aid in understanding the parallel implementations discussed in this study, this section defines key terms and MPI (message passing interface) functions that appear throughout the

algorithms discussed. MPI, message passing interface, is a standardized and portable message passing system that allows processes to run on distributed memory systems. The CPUs are physically separated from each other, but still communicate and transmit data between themselves. According to Spickler (2025), the following diagram displays how distributed memory systems look:




Figure 1: Data Flow in OpenMPI Cluster

An MPI process is a single instance of a running program in an MPI application. Each process has a unique rank, which is an integer associated with that process, separating it from the other processes. Rank 0 is typically the root process, the one that initializes the program. `MPI_Bcast` is a broadcast function that is used to send data from one process, usually the root, to all other processes in a communicator. This ensures that all processes have the same updated data before continuing. Without the usage of `MPI_Bcast`, the data would now be shared across all nodes. The following diagram displays how `MPI_Bcast` distributes memory:




Figure 2: Data Flow using MPI\_Bcast

Additionally, MPI\_Barrier is a synchronization function that blocks all processes until every process in the communicator has reached the barrier. This makes sure that all previous operations are complete across all processes before proceeding. While this tool is necessary, large overhead can occur because it pauses all processes until every one is complete. In large matrices, this build up can accumulate over time. The following diagram demonstrates how MPI\_Barrier waits for all processes to complete:

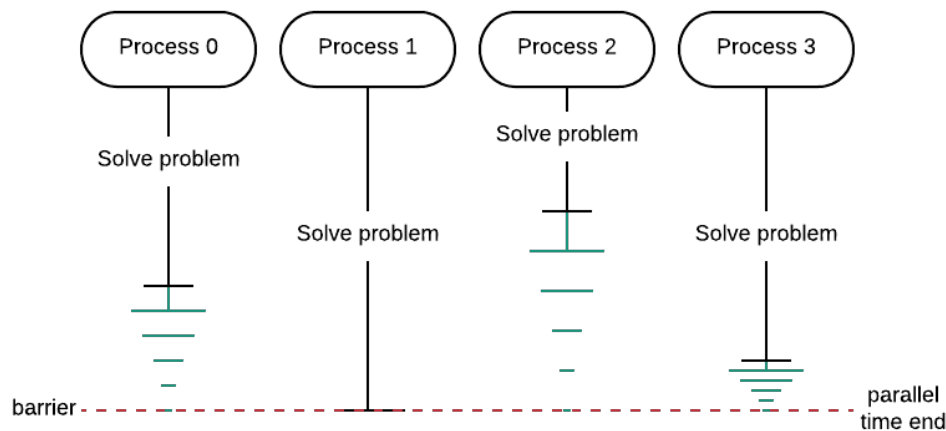


Figure 3: Data Flow using MPI\_Barrier

MPI\_Allreduce is a collective communication operation that sums values from all running processes using a specified operation, and distributes the result back to all processes. This command is used to pull all results together and re-send data back out to all processes. Here is a diagram displaying how MPI\_Allreduce works:

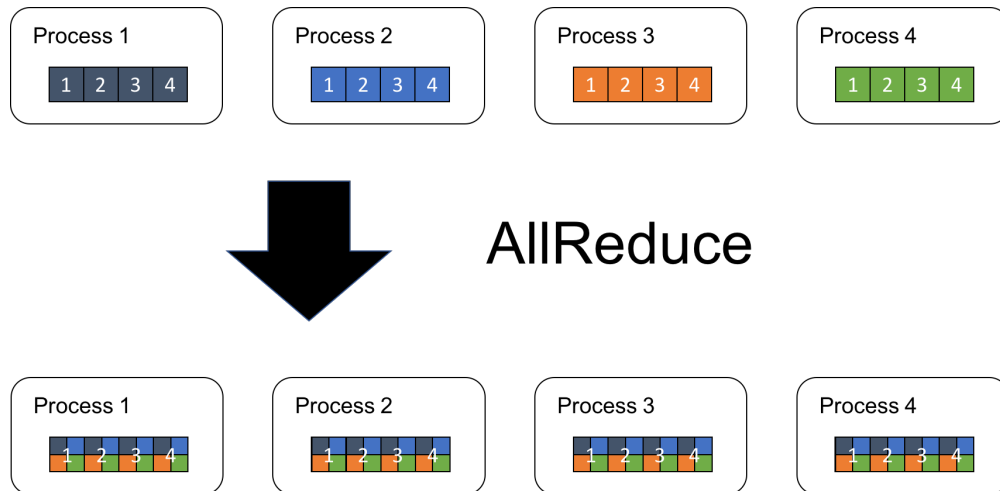


Figure 4: Data Flow using MPI\_Allreduce

### 3 Mathematics and Algorithms

The implementation of parallelized linear algebra routines requires not only a theoretical understanding of the algorithms but also practical insight into how these operations work under distributed computation. This section outlines the core mathematical algorithms used in the project: matrix multiplication, REF, RREF, and LU decomposition, as well as comparing the differences between their sequential and OpenMPI forms. Matrix multiplication is a foundational operation in linear algebra with applications ranging from linear transformations to machine learning. Given two matrices,  $A$  of size  $m \times n$  and  $B$  of size  $n \times p$ , the produce  $C = A \cdot B$  is defined as:

$$C_{i,j} = \sum_{k=1}^n A_{i,k} \cdot B_{k,j}$$

In the sequential implementation, three nested loops are used to iterate through each row  $i$  of  $A$ , each column  $j$  of  $B$ , and each inner dimension  $k$ , summing the partial products in  $C_{i,j}$ . The parallel implementation distributes the row range of matrix  $A$  across all MPI processes. Each process computes a subset of the output matrix  $C$ , then gets called by MPI\_Allreduce to consolidate the partial results across all nodes. Load balancing is managed by evenly dividing rows, and the implementations accounts for division with remainders by assigning an extra row to the first few ranks. This design allows each core to perform independent multiplication on the assigned rows, significantly reducing computation time while leveraging all available nodes. The correctness of results is guaranteed by the summation of MPI\_Allreduce.

Additionally, the REF form of a matrix is another fundamental transformation in linear algebra, and is commonly used for solving systems of linear equations and preparing a matrix for further operations such as back substitution, RREF conversion, and LU decomposition. A matrix is said to be in row echelon form when all nonzero rows are above any rows of all zeroes; the leading coefficient (pivot) of a nonzero row is strictly to the right of the pivot in the row above it; and all entries below a pivot are zero. The sequential algorithm follows the

Gaussian elimination process. The function iteratively scans for nonzero pivots, performs row swaps if needed, normalizes the pivot row so the leading entry becomes 1, and eliminates all entries below the pivot by subtracting a multiple of the pivot row from each target row. In the parallel implementation, the primary logic is preserved, but the elimination step is distributed across multiple MPI processes. Each process is responsible for updating a subset of rows. More specifically, rows are assigned based on their index modulus the total number of MPI processes:

If  $i \bmod p = \text{rank}$ , then process rank handles row  $i$ .

Each process waits until the pivot row is broadcast from rank 0, then applies the elimination step locally to its assigned rows. After completing a round of elimination, processes synchronize using `MPI_Barrier` and then broadcast the updated matrix using `MPI_Bcast` to ensure all nodes have a consistent view before moving to the next pivot row. This approach preserves correctness while achieving partial parallelism. Although some overhead exists due to the global communication at each iteration, the performance gains from dividing the row work outweigh the costs, especially on large matrices. Building upon the REF, RREF represents a stricter standard that fully solves a system of linear equations in a matrix form. In addition to the requirements for REF, RREF imposes two additional conditions: Each leading '1' is the only nonzero entry in its column; and each leading '1' is the first nonzero entry in its row. The RREF algorithm is typically used as a second phase after REF. Once the matrix is in upper triangular form, the algorithm proceeds in reverse, starting from the bottom row and moving upwards, eliminating entries above each pivot. In sequential form, the implementation iterates over the rows in reverse order. For each row, it locates the pivot, then eliminates all entries above that pivot by subtracting multiples of the pivot row from each row above. This back substitution stage ensures that each pivot is isolated. In the parallel implementation using OpenMPI, the reverse pass is distributed similarly to REF. After converting the matrix to REF form, the backward elimination steps assigns work to MPI processes based on row indices. Each process eliminates entries above pivots in rows it is responsible for, using the modulus rank pattern seen before. Once a set of updates is applied, all processes synchronize with `MPI_Barrier`, and the updated matrix is broadcasted again using `MPI_Bcast` to maintain consistency. While this introduces communication overhead, particularly in the upper triangular back propagation phase, the division of labor across processors still yields speedup in large matrices.

Furthermore, LU decomposition is another rudimentary matrix operation used to repeatedly solve systems of linear equations, especially when the coefficient matrix remains constant. The process factors a square matrix,  $A$ , into two triangular matrices— a lower triangular matrix  $L$  and an upper triangular matrix  $U$  such that  $A = L * U$ . In the sequential algorithm, Gaussian elimination is used to progressively zero out entries below the pivots, simultaneously creating the  $L$  matrix with the multipliers used during the elimination process and assigning the resulting upper triangle to  $U$ . This involves searching through each pivot position, calculating scalars for the rows below, and updating both matrices in place. In the parallel OpenMPI version, the decomposition is distributed across multiple processes by partitioning the matrix rows and assigning each subset to a different rank. During each iteration, the pivot row is broadcast from the root process to all other using `MPI_Bcast`,

and each process computes updates for its assigned rows independently. As with REF and RREF, synchronization is required after each iteration using `MPI_Barrier` to ensure consistent progress. The challenge in the parallel LU decomposition lies in how the managing data dependencies between the L and U matrices are handled.

## 4 Results

To evaluate the performance of sequential versus parallel implementations, a series of timing benchmarks were conducted on square matrices of varying sizes. Each algorithm—matrix multiplication, REF, RREF, and LU decomposition—was tested using both a single process mode and a multi processes OpenMPI mode on the constructed cluster. Matrix sized ranged from 10 x 10 to 2,000 x 2,000 and execution times were recorded for each operation. The results revealed that parallel matrix multiplication, REF, and RREF achieved significant speedups compared to their sequential counterparts, particularly as matrix size increased. For example, on a REF test of a 2,000 x 2,000 matrix, the sequential algorithm took 9.422 seconds, whereas the parallel process with 8 processes took 1.971 seconds to compute, which is roughly a 4.75x speedup. Similarly, the parallel RREF and matrix multiplication implementations showed substantial time reductions on larger matrices, although some communication overhead was observed during pivot synchronization. However, the parallel LU decomposition algorithm was almost consistently slower than the sequential version across all tested matrix sizes. For example, on a 2,000 x 2,000 matrix, the sequential time for LU Decomposition took 8.153 seconds, whereas the parallel time with 8 processes took 19.306 seconds. As discussed earlier, the method in which the algorithm has to be built and the frequent use of `MPI_Bcast` and `MPI_Barrier` caused the communication and synchronization overhead to dominate execution time, negating the benefits of multiprocessing.

Below are the timing results of all four operations being tested on the cluster in the HPCL. Variables that could affect the speeds were minimized—for example, there were no background processes running during any of the runtime tests. As seen, the sequential algorithms run slower in matrix multiplication, REF, and RREF compared to the parallel versions, whereas the sequential algorithm for LU Decomposition runs faster in comparison to the parallel version.




Figure 5: Matrix Multiplication Timings




Figure 6: REF Timings






Figure 7: RREF Timings




Figure 8: LU Decomposition Timings

## 5 Conclusion

The study demonstrated that parallel computing using OpenMPI can significantly accelerate linear algebra operations that naturally allow independent work distribution, such as matrix multiplication, REF, and RREF. Timing benchmarks confirmed that parallel implementations of these algorithms achieved substantial speedup compared to sequential versions, particularly as matrix sizes increased. However, not all algorithms benefited equally from parallelization. LU decomposition, which has strong data dependencies between elimination

steps, experienced performance degradation in the parallel implementation due to communication and synchronization overhead. This highlighted a key limitation in naive parallelization strategies—operations that frequently require inter process communication may not scale well without more sophisticated techniques or algorithms. Overall, this project provided valuable insights into both the potential and challenges of parallel computing. Future work could explore more advanced decomposition methods, such as block LU decomposition, or hybrid models that combine MPI with multithreading to better balance computation and communication.

## References

- [1] Barney, B., & Frederick, D. (n.d.). *Introduction to parallel computing tutorial*. Lawrence Livermore National Laboratory. Retrieved April 21, 2025, from <https://hpc.llnl.gov/documentation/tutorials/introduction-parallel-computing-tutorial>
- [2] Lay, D. C., Lay, S. R., & McDonald, J. J. (2020). *Linear algebra and its applications* (6th ed.). Pearson.
- [3] Spickler, D. (2025). *COSC 420: High-Performance Computing*. Salisbury University.

# MatrixC.c

```

    =1
    1
1  #include "MatrixC.h"
2
3  /*
4  Name: error()
5  Parameters: const char*
6  Return: void
7  Description: Prints error message and exits program
8  */
9  void error(const char *message)
10 {
11     fprintf(stderr, "Error: %s\n", message);
12     exit(EXIT_FAILURE);
13 }
14
15 /*
16 Name: isValid()
17 Parameters: const Matrix*
18 Return: bool
19 Description: Returns true if cols and rows are greater than 0
20 */
21 bool isValid(const Matrix *m)
22 {
23     return (m->rows > 0 && m->cols > 0);
24 }
25
26 /*
27 Name: getRows()
28 Parameters: const Matrix*
29 Return: int
30 Description: Returns the number of rows in a matrix
31 */
32 int getRows(const Matrix *m) { return m->rows; }
33
34 /*
35 Name: getCols()
36 Parameters: const Matrix*
37 Return: int
38 Description: Returns the number of cols in a matrix
39 */
40 int getCols(const Matrix *m) { return m->cols; }
41

```

```
42 /*
43 Name: init()
44 Parameters: Matrix*
45 Return: void
46 Description: Default constructor
47 */
48 void init(Matrix *m)
49 {
50     m->rows = 0;
51     m->cols = 0;
52 }
53
54 /*
55 Name: initSize()
56 Parameters: Matrix*, const int r, const int c
57 Return: void
58 Description: Constructor with set size parameters
59 */
60 void initSize(Matrix *m, const int r, const int c)
61 {
62     m->rows = r;
63     m->cols = c;
64
65     for (int i = 0; i < r; i++)
66     {
67         for (int j = 0; j < c; j++)
68         {
69             m->matrix[i][j] = 0;
70         }
71     }
72 }
73
74 /*
75 Name: initValue();
76 Parameters: Matrix*, const int r, const int c, const double
77             defval
78 Return: void
79 Description: Constructor with set size parameters and default
80             value
81 */
82 void initValue(Matrix *m, const int r, const int c, const double
83               defval)
84 {
85     m->rows = r;
86     m->cols = c;
```

```
84
85     for (int i = 0; i < r; i++)
86     {
87         for (int j = 0; j < c; j++)
88         {
89             m->matrix[i][j] = defval;
90         }
91     }
92 }
93
94 /*
95 Name: copyMatrix()
96 Parameters: Matrix *dest, const Matrix *src
97 Return: void
98 Description: Copies src matrix into dest matrix
99 */
100 void copyMatrix(Matrix *dest, const Matrix *src)
101 {
102     if (dest == src)
103         return;
104
105     if (!isValid(dest) || !isValid(src))
106     {
107         error("Cannot copy invalid matrix");
108     }
109
110     dest->rows = src->rows;
111     dest->cols = src->cols;
112
113     for (int i = 0; i < src->rows; i++)
114     {
115         for (int j = 0; j < src->cols; j++)
116         {
117             dest->matrix[i][j] = src->matrix[i][j];
118         }
119     }
120 }
121
122 /*
123 Name: display()
124 Parameters: Matrix
125 Return: void
126 Description: Displays the matrix
127 */
128 void display(const Matrix *m)
```

```

129 {
130     int maxRows = 10, maxCols = 10;
131     int showAllRows = m->rows <= maxRows;
132     int showAllCols = m->cols <= maxCols;
133
134     for (int i = 0; i < (showAllRows ? m->rows : maxRows); i++)
135     {
136         printf("[");
137         for (int j = 0; j < (showAllCols ? m->cols : maxCols); j++)
138         {
139             printf(" %7.2f", m->matrix[i][j]);
140         }
141         if (!showAllCols)
142             printf(" ...");
143         printf(" ]\n");
144     }
145     if (!showAllRows)
146     {
147         int width = (showAllCols ? m->cols : maxCols) * 8 + (
148             showAllCols ? 2 : 6);
149         for (int i = 0; i < width; i++)
150             printf(" ");
151         printf("...\n");
152     }
153
154     /*
155     Name: rowScale()
156     Parameters: Matrix*, int row, double scalar
157     Return: void
158     Description: Multiplies a row by a non-zero scalar
159     */
160     void rowScale(Matrix *m, const int row, const double scalar)
161     {
162         if (!isValid(m) || row >= m->rows || row < 0)
163         {
164             error("Invalid matrix size for row scaling");
165         }
166
167         for (int i = 0; i < m->cols; i++)
168         {
169             m->matrix[row][i] *= scalar;
170         }
171     }
172

```

```
173 /*
174 Name: rowSwap()
175 Parameters: Matrix*, int row1, int row2
176 Return: void
177 Description: Swaps two valid rows of the matrix
178 */
179 void rowSwap(Matrix *m, const int row1, const int row2)
180 {
181     if (!isValid(m) || row1 >= m->rows || row2 >= m->rows || row1 <
        0 || row2 < 0)
182     {
183         error("Invalid row access for row swap");
184     }
185
186     for (int i = 0; i < m->cols; i++)
187     {
188         double temp = m->matrix[row1][i];
189         m->matrix[row1][i] = m->matrix[row2][i];
190         m->matrix[row2][i] = temp;
191     }
192 }
193
194 /*
195 Name: rowReplace()
196 Parameters: Matrix*, int targetRow, int sourceRow, double scalar
197 Return: void
198 Description: Replaces target row with itself plus scalar
        multiplied by source row
199 */
200 void rowReplace(Matrix *m, const int targetRow, const int
        sourceRow, const double scalar)
201 {
202     if (!isValid(m) || targetRow >= m->rows || sourceRow >= m->rows
        || targetRow < 0 || sourceRow < 0)
203     {
204         error("Invalid row access for row replace");
205     }
206
207     for (int i = 0; i < m->cols; i++)
208     {
209         m->matrix[targetRow][i] += scalar * m->matrix[sourceRow][i];
210     }
211 }
212
213 /*
```



```
214 Name: setValue()
215 Parameters: Matrix*, int r, int c, double value
216 Return: void
217 Description: Sets an element of the row to a value
218 */
219 void setValue(Matrix *m, const int r, const int c, const double
    value)
220 {
221     if (!isValid(m) || r >= m->rows || c >= m->cols || r < 0 || c <
        0)
222     {
223         error("Invalid row or column access for set value");
224     }
225
226     m->matrix[r][c] = value;
227 }
228
229 /*
230 Name: setRandom()
231 Parameters: Matrix*, const int maxRand
232 Return: void
233 Description: Sets matrix to random values
234 */
235 void setRandom(Matrix *m, const int maxRand)
236 {
237     if (!isValid(m))
238     {
239         error("Invalid matrix in set random");
240     }
241
242     for (int i = 0; i < m->rows; i++)
243     {
244         for (int j = 0; j < m->cols; j++)
245         {
246             m->matrix[i][j] = rand() % maxRand;
247         }
248     }
249 }
250
251 /*
252 Name: getValue()
253 Parameters: Matrix*, const int r, const int c
254 Return: double
255 Description: Returns matrix element at a slot
256 */
```

```
257 double getValue(Matrix *m, const int r, const int c)
258 {
259     if (!isValid(m) || r >= m->rows || c >= m->cols || r < 0 || c <
        0)
260     {
261         error("Invalid row or column access in get value");
262     }
263
264     return m->matrix[r][c];
265 }
266
267 /*
268 Name: displayDimensions()
269 Parameters: Matrix*
270 Return: void
271 Description: Displays dimensions of the matrix
272 */
273 void displayDimensions(const Matrix *m)
274 {
275     if (!isValid(m))
276     {
277         error("Invalid matrix");
278     }
279
280     printf("[%d rows, %d cols]\n", m->rows, m->cols);
281 }
282
283 /*
284 Name: isSquare()
285 Parameters: const Matrix* A
286 Return: bool
287 Description: Returns true if the matrix is square
288 */
289 bool isSquare(const Matrix *A)
290 {
291     return A->rows == A->cols;
292 }
293
294 /*
295 Name: transpose()
296 Parameters: const Matrix *input, Matrix *output
297 Return: void
298 Description: Transposes the input matrix, stores result in output
        matrix
299 */
```

```
300 void transpose(const Matrix *input, Matrix *output)
301 {
302
303     if (!isValid(input))
304     {
305         error("Invalid matrix in transpose");
306     }
307
308     int rank, size;
309     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
310     MPI_Comm_size(MPI_COMM_WORLD, &size);
311     MPI_Bcast((void *)input, sizeof(Matrix), MPI_BYTE, 0,
312              MPI_COMM_WORLD);
313
314     output->rows = input->cols;
315     output->cols = input->rows;
316
317     int rows_per_proc = input->rows / size;
318     int remaining_rows = input->rows % size;
319     int start = rank * rows_per_proc + (rank < remaining_rows ?
320         rank : remaining_rows);
321     int end = start + rows_per_proc + (rank < remaining_rows ? 1 :
322         0);
323
324     for (int i = start; i < end; i++)
325     {
326         for (int j = 0; j < input->cols; j++)
327         {
328             output->matrix[j][i] = input->matrix[i][j];
329         }
330     }
331
332     // back to all processes
333     for (int proc = 0; proc < size; proc++)
334     {
335         int proc_start = proc * rows_per_proc + (proc <
336             remaining_rows ? proc : remaining_rows);
337         int proc_end = proc_start + rows_per_proc + (proc <
338             remaining_rows ? 1 : 0);
339
340         for (int row = 0; row < input->cols; row++)
341         {
342             for (int col = proc_start; col < proc_end; col++)
343             {
344                 MPI_Bcast(&output->matrix[row][col], 1, MPI_DOUBLE, proc,
```

```

        MPI_COMM_WORLD);
340     }
341 }
342 }
343 }
344
345 /*
346 Name: writeToFile()
347 Parameters: const Matrix*, const char* filename
348 Return: void
349 Description: Writes a matrix to a file
350 */
351 void writeToFile(const Matrix *m, const char *filename)
352 {
353     if (!isValid(m))
354     {
355         error("Invalid matrix");
356     }
357
358     FILE *file = fopen(filename, "w");
359     if (file == NULL)
360     {
361         error("Unable to open file for writing");
362     }
363
364     fprintf(file, "Rows: %d, Cols: %d\n", m->rows, m->cols);
365     for (int i = 0; i < 25; i++)
366         fprintf(file, "-");
367     fprintf(file, "\n");
368
369     for (int i = 0; i < m->rows; i++)
370     {
371         for (int j = 0; j < m->cols; j++)
372         {
373             fprintf(file, "%lf ", m->matrix[i][j]);
374         }
375         fprintf(file, "\n");
376     }
377
378     fclose(file);
379 }
380
381 /*
382 Name: solve()
383 Parameters: const Matrix*, const double[], const int size

```

```

384 Return: bool
385 Description: Returns true if the solution set works with the
      matrix
386 */
387 bool solve(const Matrix *m, const double solutions[], const int
      size)
388 {
389     if (!isValid(m) || size < 1)
390     {
391         error("Cannot check solution set on empty matrix or empty
            solution set");
392     }
393
394     const double EPSILON = 1e-5;
395     double value = 0;
396
397     for (int i = 0; i < m->rows; i++)
398     {
399         value = 0;
400         for (int j = 0; j < m->cols - 1; j++)
401         {
402             value += m->matrix[i][j] * solutions[j];
403         }
404         if (fabs(value - m->matrix[i][m->cols - 1] > EPSILON))
405             return false;
406     }
407     return true;
408 }
409
410 /*
411 Name: multiplyMatrix()
412 Parameters: const Matrix *A, const Matrix *B, Matrix *C
413 Return: void
414 Description: Multiplies matrix A and B using OpenMPI and stores
      result in matrix C.
415 */
416 void multiplyMatrix(const Matrix *A, const Matrix *B, Matrix *C)
417 {
418     int rank, size;
419     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
420     MPI_Comm_size(MPI_COMM_WORLD, &size);
421
422     int a_rows = A->rows;
423     int a_cols = A->cols;
424     int b_cols = B->cols;

```

```

425
426     if (rank == 0)
427     {
428         if (!isValid(A) || !isValid(B))
429             error("Invalid input matrices");
430         if (a_cols != B->rows)
431             error("Incompatible dimensions for multiplication");
432     }
433
434     // Broadcast matrix dimensions
435     MPI_Bcast(&a_rows, 1, MPI_INT, 0, MPI_COMM_WORLD);
436     MPI_Bcast(&a_cols, 1, MPI_INT, 0, MPI_COMM_WORLD);
437     MPI_Bcast(&b_cols, 1, MPI_INT, 0, MPI_COMM_WORLD);
438
439     // Broadcast B (whole matrix)
440     MPI_Bcast((void *)B->matrix, MAX_ROWS * MAX_COLS, MPI_DOUBLE,
441               0, MPI_COMM_WORLD);
442
443     // Calculate local rows
444     int rows_per_proc = a_rows / size;
445     int remainder = a_rows % size;
446     int local_rows = (rank < remainder) ? rows_per_proc + 1 :
447                     rows_per_proc;
448     printf("Rank %d received %d rows, computing...\n", rank,
449           local_rows);
450
451     // Allocate only as much as needed
452     double local_A[local_rows][a_cols];
453     double local_C[local_rows][b_cols];
454     memset(local_C, 0, sizeof(local_C));
455
456     // Prepare scatter metadata
457     int sendcounts[size];
458     int displs[size];
459     int offset = 0;
460     for (int i = 0; i < size; i++)
461     {
462         int count = ((i < remainder) ? rows_per_proc + 1 :
463                     rows_per_proc) * a_cols;
464         sendcounts[i] = count;
465         displs[i] = offset;
466         offset += count;
467     }
468
469     // Scatter A rows to processes

```

```

466 MPI_Scatterv(&(A->matrix[0][0]), sendcounts, displs, MPI_DOUBLE
467             ,
468             &(local_A[0][0]), local_rows * a_cols, MPI_DOUBLE,
469             0, MPI_COMM_WORLD);
470 // Compute local result
471 for (int i = 0; i < local_rows; i++)
472 {
473     for (int j = 0; j < b_cols; j++)
474     {
475         for (int k = 0; k < a_cols; k++)
476         {
477             local_C[i][j] += local_A[i][k] * B->matrix[k][j];
478         }
479     }
480 }
481 // Prepare gather metadata
482 int recvcounts[size];
483 int recvdispls[size];
484 offset = 0;
485 for (int i = 0; i < size; i++)
486 {
487     int count = ((i < remainder) ? rows_per_proc + 1 :
488                 rows_per_proc) * b_cols;
489     recvcounts[i] = count;
490     recvdispls[i] = offset;
491     offset += count;
492 }
493 // Gather results into final matrix C
494 MPI_Gatherv(&(local_C[0][0]), local_rows * b_cols, MPI_DOUBLE,
495            &(C->matrix[0][0]), recvcounts, recvdispls,
496            MPI_DOUBLE,
497            0, MPI_COMM_WORLD);
498 // Set dimensions on root
499 if (rank == 0)
500 {
501     C->rows = a_rows;
502     C->cols = b_cols;
503 }
504 }
505 }
506
507 /*

```

```

508 Name: ref()
509 Parameters: Matrix *m
510 Return: void
511 Description: Computes the Row Echelon Form (REF) of matrix m
               using OpenMPI.
512 */
513 void ref(Matrix *m)
514 {
515     if (!isValid(m))
516         error("Invalid matrix");
517
518     int rank, size;
519     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
520     MPI_Comm_size(MPI_COMM_WORLD, &size);
521
522     int rows = m->rows;
523     int cols = m->cols;
524
525     int rows_per_proc = rows / size;
526     int remainder = rows % size;
527     int local_rows = (rank < remainder) ? rows_per_proc + 1 :
                       rows_per_proc;
528     printf("Rank %d preparing to allocate %d rows    %d cols = %.2f
           MB\n", rank, local_rows, cols, (local_rows * cols * sizeof(
           double)) / (1024.0 * 1024.0));
529
530     int start_row = (rank < remainder)
531                     ? rank * (rows_per_proc + 1)
532                     : rank * rows_per_proc + remainder;
533
534     double local_matrix[local_rows][cols];
535     printf("Rank %d received %d rows, computing...\n", rank,
           local_rows);
536
537     // Scatter rows of m into local_matrix
538     int sendcounts[size], displs[size];
539     int offset = 0;
540     for (int i = 0; i < size; i++)
541     {
542         int count = ((i < remainder) ? rows_per_proc + 1 :
           rows_per_proc) * cols;
543         sendcounts[i] = count;
544         displs[i] = offset;
545         offset += count;
546     }

```



```

547
548 MPI_Scatterv(&(m->matrix[0][0]), sendcounts, displs, MPI_DOUBLE
549             ,
550             &(local_matrix[0][0]), local_rows * cols,
551             MPI_DOUBLE,
552             0, MPI_COMM_WORLD);
553
554 // Temporary pivot row buffer
555 double pivot_row[cols];
556
557 for (int r = 0; r < rows; r++)
558 {
559     int owner = -1, local_r = -1;
560     for (int i = 0; i < size; i++)
561     {
562         int start = (i < remainder) ? i * (rows_per_proc + 1) : i *
563             rows_per_proc + remainder;
564         int end = start + ((i < remainder) ? rows_per_proc + 1 :
565             rows_per_proc);
566         if (r >= start && r < end)
567         {
568             owner = i;
569             local_r = r - start;
570             break;
571         }
572     }
573
574 // Broadcast pivot row
575 if (rank == owner)
576     memcpy(pivot_row, local_matrix[local_r], sizeof(double) *
577         cols);
578
579 MPI_Bcast(pivot_row, cols, MPI_DOUBLE, owner, MPI_COMM_WORLD)
580 ;
581
582 // Eliminate below pivot
583 for (int i = 0; i < local_rows; i++)
584 {
585     int global_row = start_row + i;
586     if (global_row <= r)
587         continue;
588
589     double factor = local_matrix[i][r] / pivot_row[r];
590     for (int j = r; j < cols; j++)
591     {

```

```

586         local_matrix[i][j] -= factor * pivot_row[j];
587     }
588 }
589 }
590
591 // Gather results back to m
592 MPI_Gatherv(&(local_matrix[0][0]), local_rows * cols,
593            MPI_DOUBLE,
594            &(m->matrix[0][0]), sendcounts, displs, MPI_DOUBLE,
595            0, MPI_COMM_WORLD);
596 }
597 /*
598 Name: rref()
599 Parameters: Matrix *m
600 Return: void
601 Description: Computes Reduced Row Echelon Form (RREF) using
602             OpenMPI.
603 */
604 void rref(Matrix *m)
605 {
606     if (!isValid(m))
607         error("Invalid matrix");
608
609     int rank, size;
610     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
611     MPI_Comm_size(MPI_COMM_WORLD, &size);
612
613     int rows = m->rows;
614     int cols = m->cols;
615
616     int rows_per_proc = rows / size;
617     int remainder = rows % size;
618     int local_rows = (rank < remainder) ? rows_per_proc + 1 :
619                     rows_per_proc;
620
621     int start_row = (rank < remainder)
622                     ? rank * (rows_per_proc + 1)
623                     : rank * rows_per_proc + remainder;
624
625     double local_matrix[local_rows][cols];
626
627     printf("Rank %d received %d rows, computing...\n", rank,
628           local_rows);

```

```

627 // Setup for scatter/gather
628 int sendcounts[size], displs[size];
629 int offset = 0;
630 for (int i = 0; i < size; i++)
631 {
632     int count = ((i < remainder) ? rows_per_proc + 1 :
633                 rows_per_proc) * cols;
634     sendcounts[i] = count;
635     displs[i] = offset;
636     offset += count;
637 }
638 // Scatter matrix
639 MPI_Scatterv(&(m->matrix[0][0]), sendcounts, displs, MPI_DOUBLE
640             ,
641             &(local_matrix[0][0]), local_rows * cols,
642             MPI_DOUBLE,
643             0, MPI_COMM_WORLD);
644
645 double pivot_row[cols];
646
647 for (int r = 0; r < rows; r++)
648 {
649     int owner = -1, local_r = -1;
650     for (int i = 0; i < size; i++)
651     {
652         int s = (i < remainder) ? i * (rows_per_proc + 1) : i *
653             rows_per_proc + remainder;
654         int e = s + ((i < remainder) ? rows_per_proc + 1 :
655             rows_per_proc);
656         if (r >= s && r < e)
657         {
658             owner = i;
659             local_r = r - s;
660             break;
661         }
662     }
663
664 // Normalize pivot row on owner
665 if (rank == owner)
666 {
667     double pivot = local_matrix[local_r][r];
668     if (pivot != 0)
669     {
670         for (int j = r; j < cols; j++)

```

```

667         local_matrix[local_r][j] /= pivot;
668     }
669     memcpy(pivot_row, local_matrix[local_r], sizeof(double) *
        cols);
670 }
671
672 // Broadcast normalized pivot row
673 MPI_Bcast(pivot_row, cols, MPI_DOUBLE, owner, MPI_COMM_WORLD)
        ;
674
675 // Eliminate all other rows (above and below)
676 for (int i = 0; i < local_rows; i++)
677 {
678     int global_i = start_row + i;
679     if (global_i == r)
680         continue;
681
682     double factor = local_matrix[i][r];
683     for (int j = r; j < cols; j++)
684         local_matrix[i][j] -= factor * pivot_row[j];
685 }
686 }
687
688 // Gather result to root
689 MPI_Gatherv(&(local_matrix[0][0]), local_rows * cols,
        MPI_DOUBLE,
690             &(m->matrix[0][0]), sendcounts, displs, MPI_DOUBLE,
691             0, MPI_COMM_WORLD);
692 }
693
694 /*
695 Name: addMatrix()
696 Parameters: const Matrix *A, const Matrix *B, Matrix *result
697 Return: Void
698 Description: Stores result of addition of A and B
699 */
700 void addMatrix(const Matrix *A, const Matrix *B, Matrix *result)
701 {
702     if (!isValid(A) || !isValid(B) || A->cols != B->cols || A->rows
        != B->rows)
703     {
704         error("Invalid matrix size");
705     }
706
707     int rank, size;

```

```

708 MPI_Comm_rank(MPI_COMM_WORLD, &rank);
709 MPI_Comm_size(MPI_COMM_WORLD, &size);
710
711 int rows_per_proc = A->rows / size;
712 int remaining_rows = A->rows % size;
713 int start = rank * rows_per_proc + (rank < remaining_rows ?
    rank : remaining_rows);
714 int end = start + rows_per_proc + (rank < remaining_rows ? 1 :
    0);
715
716 for (int i = start; i < end; i++)
717 {
718     for (int j = 0; j < A->cols; j++)
719     {
720         result->matrix[i][j] = A->matrix[i][j] + B->matrix[i][j];
721     }
722 }
723
724 MPI_Barrier(MPI_COMM_WORLD);
725 MPI_Allreduce(MPI_IN_PLACE, result->matrix, MAX_ROWS * MAX_COLS
    , MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
726 }
727
728 /*
729 Name: subtractMatrix()
730 Parameters: const Matrix *A, const Matrix *B, Matrix *result
731 Return: Void
732 Description: Stores result of subtraction of A and B
733 */
734 void subtractMatrix(const Matrix *A, const Matrix *B, Matrix *
    result)
735 {
736     if (!isValid(A) || !isValid(B) || A->cols != B->cols || A->rows
        != B->rows)
737     {
738         error("Invalid matrix size");
739     }
740
741     if (!isValid(A) || !isValid(B) || A->cols != B->cols || A->rows
        != B->rows)
742     {
743         error("Invalid matrix size");
744     }
745
746     int rank, size;

```

```

747 MPI_Comm_rank(MPI_COMM_WORLD, &rank);
748 MPI_Comm_size(MPI_COMM_WORLD, &size);
749
750 int rows_per_proc = A->rows / size;
751 int remaining_rows = A->rows % size;
752 int start = rank * rows_per_proc + (rank < remaining_rows ?
    rank : remaining_rows);
753 int end = start + rows_per_proc + (rank < remaining_rows ? 1 :
    0);
754
755 for (int i = start; i < end; i++)
756 {
757     for (int j = 0; j < A->cols; j++)
758     {
759         result->matrix[i][j] = A->matrix[i][j] - B->matrix[i][j];
760     }
761 }
762
763 MPI_Barrier(MPI_COMM_WORLD);
764 MPI_Allreduce(MPI_IN_PLACE, result->matrix, MAX_ROWS * MAX_COLS
    , MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
765 }
766
767 /*
768 Name: LU
769 Parameters: Matrix* A, Matrix* L, Matrix* U
770 Return: void
771 Description: Performs LU decomposition
772 */
773 void LU(Matrix *A, Matrix *L, Matrix *U)
774 {
775     int n = A->rows;
776
777     // Initialize L and U matrices
778     for (int i = 0; i < n; ++i)
779     {
780         for (int j = 0; j < n; ++j)
781         {
782             L->matrix[i][j] = 0.0;
783             U->matrix[i][j] = A->matrix[i][j];
784         }
785     }
786
787     for (int k = 0; k < n; ++k)
788     {

```

```

789     for (int i = k + 1; i < n; ++i)
790     {
791         double multiplier = U->matrix[i][k] / U->matrix[k][k];
792         L->matrix[i][k] = multiplier;
793
794         for (int j = k; j < n; ++j)
795         {
796             U->matrix[i][j] -= multiplier * U->matrix[k][j];
797         }
798     }
799 }
800
801 // Set the diagonal of L to 1
802 for (int i = 0; i < n; ++i)
803 {
804     L->matrix[i][i] = 1.0;
805 }
806 }
807
808 /*
809
810
811
812
813
814 Sequential functions below
815
816
817
818
819
820 */
821
822 /*
823 Name: Seq_multiplyMatrix()
824 Parameters: const Matrix *A, const Matrix *B, Matrix *C
825 Return: void
826 Description: Multiplies matrix A and matrix B sequentially and
827               stores result in matrix C.
828 */
829 void Seq_multiplyMatrix(const Matrix *A, const Matrix *B, Matrix
830                        *C)
831 {
832     if (!isValid(A) || !isValid(B))
833         error("Invalid input matrices");

```

```

832     if (A->cols != B->rows)
833         error("Incompatible dimensions for multiplication");
834
835     for (int i = 0; i < A->rows; i++)
836     {
837         for (int j = 0; j < B->cols; j++)
838         {
839             C->matrix[i][j] = 0;
840             for (int k = 0; k < A->cols; k++)
841             {
842                 C->matrix[i][j] += A->matrix[i][k] * B->matrix[k][j];
843             }
844         }
845     }
846 }
847
848 /*
849 Name: Seq_ref()
850 Parameters: Matrix *m
851 Return: void
852 Description: Puts the matrix into row echelon form sequentially.
853 */
854 void Seq_ref(Matrix *m)
855 {
856     if (!isValid(m))
857         error("Invalid matrix");
858
859     int lead = 0;
860     int rowCount = m->rows;
861     int colCount = m->cols;
862
863     for (int r = 0; r < rowCount; r++)
864     {
865         if (lead >= colCount)
866             return;
867
868         int i = r;
869         while (m->matrix[i][lead] == 0)
870         {
871             i++;
872             if (i == rowCount)
873             {
874                 i = r;
875                 lead++;
876                 if (lead == colCount)

```



```

877         return;
878     }
879 }
880
881 if (i != r)
882 {
883     for (int j = 0; j < colCount; j++)
884     {
885         double temp = m->matrix[r][j];
886         m->matrix[r][j] = m->matrix[i][j];
887         m->matrix[i][j] = temp;
888     }
889 }
890
891 double lv = m->matrix[r][lead];
892 if (lv != 0)
893 {
894     for (int j = 0; j < colCount; j++)
895         m->matrix[r][j] /= lv;
896 }
897
898 for (int i = r + 1; i < rowCount; i++)
899 {
900     double lv2 = m->matrix[i][lead];
901     for (int j = 0; j < colCount; j++)
902         m->matrix[i][j] -= lv2 * m->matrix[r][j];
903 }
904
905 lead++;
906 }
907 }
908
909 /*
910 Name: Seq_rref()
911 Parameters: Matrix *m
912 Return: void
913 Description: Converts matrix to reduced row echelon form without
914             MPI.
915 */
916 void Seq_rref(Matrix *m)
917 {
918     if (!isValid(m))
919         error("Invalid matrix");
920
921     Seq_ref(m);

```

```

921
922     int rowCount = m->rows;
923     int colCount = m->cols;
924
925     for (int r = rowCount - 1; r >= 0; r--)
926     {
927         int leadCol = -1;
928         for (int j = 0; j < colCount; j++)
929         {
930             if (fabs(m->matrix[r][j] - 1.0) < 1e-6)
931             {
932                 leadCol = j;
933                 break;
934             }
935         }
936
937         if (leadCol == -1)
938             continue;
939
940         for (int i = 0; i < r; i++)
941         {
942             double factor = m->matrix[i][leadCol];
943             for (int j = 0; j < colCount; j++)
944                 m->matrix[i][j] -= factor * m->matrix[r][j];
945         }
946     }
947 }
948
949 /*
950 Name: Seq_LU()
951 Parameters: const Matrix *A, Matrix *L, Matrix *U
952 Return: void
953 Description: Performs LU decomposition sequentially.
954 */
955 void Seq_LU(const Matrix *A, Matrix *L, Matrix *U)
956 {
957     if (!isValid(A))
958         error("Invalid matrix");
959
960     int n = A->rows;
961
962     for (int i = 0; i < n; ++i)
963     {
964         for (int j = 0; j < n; ++j)
965         {

```

```
966     L->matrix[i][j] = 0;
967     U->matrix[i][j] = 0;
968 }
969 }
970
971 for (int k = 0; k < n; ++k)
972 {
973     for (int j = k; j < n; ++j)
974     {
975         double sum = 0;
976         for (int s = 0; s < k; ++s)
977             sum += L->matrix[k][s] * U->matrix[s][j];
978         U->matrix[k][j] = A->matrix[k][j] - sum;
979     }
980
981     L->matrix[k][k] = 1.0;
982
983     for (int i = k + 1; i < n; ++i)
984     {
985         double sum = 0;
986         for (int s = 0; s < k; ++s)
987             sum += L->matrix[i][s] * U->matrix[s][k];
988         L->matrix[i][k] = (A->matrix[i][k] - sum) / U->matrix[k][k];
989     }
990 }
991 }
```

## multiplyTest.c

```

=1
1
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <mpi.h>
4 #include <time.h>
5 #include "MatrixC.h"
6
7 void divisor(const char *msg)
8 {
9     printf("\n---- %s ----\n\n", msg);
10 }
11
12 Matrix A, B, C_parallel, C_seq;
13
14 int main(int argc, char **argv)
15 {
16     MPI_Init(&argc, &argv);
17
18     int rank, size;
19     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
20     MPI_Comm_size(MPI_COMM_WORLD, &size);
21
22     srand(time(NULL) + rank);
23
24     int N = 100, use_parallel = 0;
25
26     if (rank == 0)
27     {
28         printf("Enter matrix size: ");
29         scanf("%d", &N);
30
31         printf("Run sequential too? (1 = yes, 0 = no): ");
32         scanf("%d", &use_parallel);
33     }
34
35     // Broadcast settings
36     MPI_Bcast(&N, 1, MPI_INT, 0, MPI_COMM_WORLD);
37     MPI_Bcast(&use_parallel, 1, MPI_INT, 0, MPI_COMM_WORLD);
38
39     initSize(&A, N, N);
40     initSize(&B, N, N);
41     initSize(&C_parallel, N, N);

```

```

42  initSize(&C_seq, N, N);
43
44  // Rank 0 generates the random data
45  if (rank == 0)
46  {
47      setRandom(&A, 100);
48      setRandom(&B, 100);
49  }
50
51  // Broadcast A and B
52  MPI_Bcast(&(A.matrix[0][0]), N * N, MPI_DOUBLE, 0,
53           MPI_COMM_WORLD);
54  MPI_Bcast(&(B.matrix[0][0]), N * N, MPI_DOUBLE, 0,
55           MPI_COMM_WORLD);
56
57  // ----- Parallel -----
58  if (rank == 0)
59      printf("Running parallel multiplyMatrix()...\n");
60
61  double start = MPI_Wtime();
62  multiplyMatrix(&A, &B, &C_parallel);
63  double end = MPI_Wtime();
64
65  if (rank == 0)
66  {
67      printf("multiplyMatrix() done in %.3f seconds\n", end - start);
68      divisor("multiplyMatrix");
69  }
70
71  // ----- Sequential -----
72  if (rank == 0 && use_parallel == 1)
73  {
74      printf("Running sequential Seq_multiplyMatrix()...\n");
75
76      start = MPI_Wtime();
77      Seq_multiplyMatrix(&A, &B, &C_seq);
78      end = MPI_Wtime();
79
80      printf("Seq_multiplyMatrix() done in %.3f seconds\n", end -
81             start);
82      divisor("Seq_multiplyMatrix");
83  }
84
85  MPI_Finalize();

```

```
83     return 0;  
84 }
```

## refTest.c

```

=1
1
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <mpi.h>
4 #include <time.h>
5 #include "MatrixC.h"
6
7 void divisor(const char *msg)
8 {
9     printf("\n---- %s ----\n\n", msg);
10 }
11
12 // Global matrices
13 Matrix A_parallel, A_seq;
14
15 int main(int argc, char **argv)
16 {
17     MPI_Init(&argc, &argv);
18
19     int rank, size;
20     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
21     MPI_Comm_size(MPI_COMM_WORLD, &size);
22
23     srand(time(NULL) + rank);
24
25     int N = 10, use_parallel = 1;
26
27     if (rank == 0)
28     {
29         printf("Enter matrix size: ");
30         scanf("%d", &N);
31
32         printf("Run sequential too? (1 = yes, 0 = no): ");
33         scanf("%d", &use_parallel);
34     }
35
36     // Broadcast settings
37     MPI_Bcast(&N, 1, MPI_INT, 0, MPI_COMM_WORLD);
38     MPI_Bcast(&use_parallel, 1, MPI_INT, 0, MPI_COMM_WORLD);
39
40     initSize(&A_parallel, N, N);
41     initSize(&A_seq, N, N);

```

```
42
43 // Rank 0 generates the random matrix
44 if (rank == 0)
45 {
46     setRandom(&A_parallel, 100);
47     copyMatrix(&A_seq, &A_parallel);
48 }
49
50 // Broadcast A_parallel to all ranks
51 MPI_Bcast(&(A_parallel.matrix[0][0]), N * N, MPI_DOUBLE, 0,
52          MPI_COMM_WORLD);
53
54 // ----- Parallel ref() -----
55 if (rank == 0)
56     printf("Running parallel ref()...\n");
57
58 double start = MPI_Wtime();
59 ref(&A_parallel);
60 double end = MPI_Wtime();
61
62 if (rank == 0)
63 {
64     printf("ref() done in %.3f seconds\n", end - start);
65     divisor("ref");
66 }
67
68 // ----- Sequential Seq_ref() -----
69 if (rank == 0 && use_parallel == 1)
70 {
71     printf("Running sequential Seq_ref()...\n");
72
73     start = MPI_Wtime();
74     Seq_ref(&A_seq);
75     end = MPI_Wtime();
76
77     printf("Seq_ref() done in %.3f seconds\n", end - start);
78     divisor("Seq_ref");
79 }
80
81 MPI_Finalize();
82 return 0;
83 }
```



## rrefTest.c

```

=1
1
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <mpi.h>
4 #include <time.h>
5 #include "MatrixC.h"
6
7 void divisor(const char *msg)
8 {
9     printf("\n---- %s ----\n\n", msg);
10 }
11
12 // Global matrices
13 Matrix A_parallel, A_seq;
14
15 int main(int argc, char **argv)
16 {
17     MPI_Init(&argc, &argv);
18
19     int rank, size;
20     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
21     MPI_Comm_size(MPI_COMM_WORLD, &size);
22
23     srand(time(NULL) + rank);
24
25     int N = 10, use_parallel = 1;
26
27     if (rank == 0)
28     {
29         printf("Enter matrix size: ");
30         scanf("%d", &N);
31
32         printf("Run sequential too? (1 = yes, 0 = no): ");
33         scanf("%d", &use_parallel);
34     }
35
36     // Broadcast settings
37     MPI_Bcast(&N, 1, MPI_INT, 0, MPI_COMM_WORLD);
38     MPI_Bcast(&use_parallel, 1, MPI_INT, 0, MPI_COMM_WORLD);
39
40     initSize(&A_parallel, N, N);
41     initSize(&A_seq, N, N);

```

```
42
43 // Rank 0 sets data
44 if (rank == 0)
45 {
46     setRandom(&A_parallel, 100);
47     copyMatrix(&A_seq, &A_parallel);
48 }
49
50 // Broadcast to all processes
51 MPI_Bcast(&(A_parallel.matrix[0][0]), N * N, MPI_DOUBLE, 0,
52          MPI_COMM_WORLD);
53
54 // ----- Parallel rref() -----
55 if (rank == 0)
56     printf("Running parallel rref()...\n");
57
58 double start = MPI_Wtime();
59 rref(&A_parallel);
60 double end = MPI_Wtime();
61
62 if (rank == 0)
63 {
64     printf("rref() done in %.3f seconds\n", end - start);
65     divisor("rref");
66 }
67
68 // ----- Sequential Seq_rref() -----
69 if (rank == 0 && use_parallel == 1)
70 {
71     printf("Running sequential Seq_rref()...\n");
72
73     start = MPI_Wtime();
74     Seq_rref(&A_seq);
75     end = MPI_Wtime();
76
77     printf("Seq_rref() done in %.3f seconds\n", end - start);
78     divisor("Seq_rref");
79 }
80 MPI_Finalize();
81 return 0;
82 }
```

# luTest.c

```

=1
1
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <mpi.h>
4 #include <time.h>
5 #include "MatrixC.h"
6
7 void divisor(const char *msg);
8
9 Matrix A_original, L_parallel, U_parallel;
10 Matrix L_seq, U_seq;
11
12 int main(int argc, char **argv)
13 {
14     MPI_Init(&argc, &argv);
15
16     int rank, size;
17     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
18     MPI_Comm_size(MPI_COMM_WORLD, &size);
19
20     srand(time(NULL) + rank);
21
22     int N = 10, use_sequential = 0;
23
24     if (rank == 0)
25     {
26         printf("Enter matrix size (N for NxN): ");
27         scanf("%d", &N);
28
29         printf("Run sequential also? (1 = yes, 0 = no): ");
30         scanf("%d", &use_sequential);
31     }
32
33     MPI_Bcast(&N, 1, MPI_INT, 0, MPI_COMM_WORLD);
34     MPI_Bcast(&use_sequential, 1, MPI_INT, 0, MPI_COMM_WORLD);
35
36     initSize(&A_original, N, N);
37     initSize(&L_parallel, N, N);
38     initSize(&U_parallel, N, N);
39
40     if (rank == 0)
41     {

```

```

42     initSize(&L_seq, N, N);
43     initSize(&U_seq, N, N);
44     setRandom(&A_original, 100);
45 }
46
47 MPI_Bcast(&(A_original.matrix[0][0]), N * N, MPI_DOUBLE, 0,
48          MPI_COMM_WORLD);
49
50 // ----- Parallel LU -----
51 if (rank == 0)
52     printf("Running parallel LU()...\n");
53
54 double start = MPI_Wtime();
55 LU(&A_original, &L_parallel, &U_parallel);
56 double end = MPI_Wtime();
57
58 if (rank == 0)
59 {
60     printf("Parallel LU() done in %.3f seconds\n", end - start);
61     divisor("Parallel LU");
62
63     if (use_sequential == 1)
64     {
65         Matrix A_copy;
66         initSize(&A_copy, N, N);
67         copyMatrix(&A_copy, &A_original);
68
69         printf("Running sequential Seq_LU()...\n");
70
71         start = MPI_Wtime();
72         Seq_LU(&A_copy, &L_seq, &U_seq);
73         end = MPI_Wtime();
74
75         printf("Sequential Seq_LU() done in %.3f seconds\n", end -
76              start);
77         divisor("Sequential LU");
78     }
79 }
80 MPI_Finalize();
81 return 0;
82
83 void divisor(const char *msg)
84 {

```

```
85     printf("\n---- %s ----\n\n", msg);  
86 }
```

## runTest.sh

```

=1
1
1 #!/bin/bash
2
3 echo "Enter matrix size (N for NxN): "
4 read N
5
6 echo "Run sequential also? (1 = yes, 0 = no): "
7 read USE_SEQ
8
9 echo "Enter number of MPI processes to use: "
10 read PROCS
11
12 echo "Which test to run?"
13 echo "1. multiplyTest"
14 echo "2. refTest"
15 echo "3. rrefTest"
16 echo "4. luTest"
17 read CHOICE
18
19 # Pick executable
20 case $CHOICE in
21 1) EXEC="./multiplyTest" ;;
22 2) EXEC="./refTest" ;;
23 3) EXEC="./rrefTest" ;;
24 4) EXEC="./luTest" ;;
25 *) echo "Invalid choice"; exit 1 ;;
26 esac
27
28 echo
29 echo "Running $EXEC with:"
30 echo "- Matrix size: $N"
31 echo "- Sequential: $USE_SEQ"
32 echo "- Processes: $PROCS"
33 echo
34
35 mpirun --oversubscribe -np $PROCS $EXEC <<< "$N
36 $USE_SEQ"

```